

```
In [1]: import sqlite3
import pandas as pd
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
```

1. Load the Data (sql db)

```
In [2]: con = sqlite3.connect("bmarket.db")
```

```
In [3]: cursor = con.cursor()

# Query the sqlite_master table to get table names
cursor.execute("SELECT name FROM sqlite_master WHERE type='table';")

# Fetch all results
table_names = [row[0] for row in cursor.fetchall()]
print("Tables in the database:", table_names)
```

Tables in the database: ['bank_marketing']

```
In [4]: query = "SELECT * FROM bank_marketing"
df = pd.read_sql_query(query, con)

con.close()

df.head()
```

Out[4]:

| | Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Personal Loan | Contact Method |
|---|-----------|----------|-------------|----------------|-----------------|----------------|--------------|---------------|----------------|
| 0 | 32885 | 57 years | technician | married | high.school | no | no | yes | Cellular |
| 1 | 3170 | 55 years | unknown | married | unknown | unknown | yes | no | telephone |
| 2 | 32207 | 33 years | blue-collar | married | basic.9y | no | no | no | cellular |
| 3 | 9404 | 36 years | admin. | married | high.school | no | no | no | Telephone |
| 4 | 14021 | 27 years | housemaid | married | high.school | no | None | no | Cellular |

2. Initial Data Understanding

We first need to inspect the dataset, to gain a better understanding of what we are working with.

This includes:

- Number of rows & columns
- Column names & data types
- Summary of the statistics using .describe()
- Outliers/Data quality issues

```
In [5]: df.columns
```

```
Out[5]: Index(['Client ID', 'Age', 'Occupation', 'Marital Status', 'Education Level',
   'Credit Default', 'Housing Loan', 'Personal Loan', 'Contact Method',
   'Campaign Calls', 'Previous Contact Days', 'Subscription Status'],
  dtype='object')
```

```
In [6]: df.shape
```

```
Out[6]: (41188, 12)
```

By looking at the results from "df.describe(include='all').T", we immediately make a few observations:

-Under the 'top' result for the Age feature, it shows the highest age being 150 years old. As of now, the oldest person alive is 116 years old & the oldest ever recorded was 122. We need to dive into this later and verify the information

```
In [7]: df.describe(include='all').T
```

Out[7]:

| | | count | unique | top | freq | mean | std | min |
|------------------------------|---------|-------|-------------------|-------|------|------------|-------------|-------|
| Client ID | 41188.0 | NaN | | NaN | NaN | 20594.5 | 11890.09578 | 1.0 |
| Age | 41188 | 77 | 150 years | 4197 | NaN | NaN | NaN | NaN |
| Occupation | 41188 | 12 | admin. | 10422 | NaN | NaN | NaN | NaN |
| Marital Status | 41188 | 4 | married | 24928 | NaN | NaN | NaN | NaN |
| Education Level | 41188 | 8 | university.degree | 12168 | NaN | NaN | NaN | NaN |
| Credit Default | 41188 | 3 | no | 32588 | NaN | NaN | NaN | NaN |
| Housing Loan | 16399 | 3 | yes | 8595 | NaN | NaN | NaN | NaN |
| Personal Loan | 37042 | 3 | no | 30532 | NaN | NaN | NaN | NaN |
| Contact Method | 41188 | 4 | Cell | 13100 | NaN | NaN | NaN | NaN |
| Campaign Calls | 41188.0 | NaN | | NaN | NaN | 2.051374 | 3.171345 | -41.0 |
| Previous Contact Days | 41188.0 | NaN | | NaN | NaN | 962.475454 | 186.910907 | 0.0 |
| Subscription Status | 41188 | 2 | no | 36548 | NaN | NaN | NaN | NaN |



Insights from running .duplicated().sum()

There are no duplicated entries inside the dataset, this helps to ensure reliable data that does not skew predictions

In [8]: `df.duplicated().sum()`

Out[8]: `np.int64(0)`

Insights from running .info()

There is 1 column currently stored as the 'object' type that is using the wrong data type, and needs to be changed:

-Age

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Client ID        41188 non-null   int64  
 1   Age              41188 non-null   object  
 2   Occupation       41188 non-null   object  
 3   Marital Status   41188 non-null   object  
 4   Education Level  41188 non-null   object  
 5   Credit Default   41188 non-null   object  
 6   Housing Loan     16399 non-null   object  
 7   Personal Loan    37042 non-null   object  
 8   Contact Method   41188 non-null   object  
 9   Campaign Calls   41188 non-null   int64  
 10  Previous Contact Days 41188 non-null   int64  
 11  Subscription Status 41188 non-null   object  
dtypes: int64(3), object(9)
memory usage: 3.8+ MB
```

```
In [10]: df.isna().sum()
```

```
Out[10]: Client ID          0
Age                0
Occupation        0
Marital Status    0
Education Level   0
Credit Default    0
Housing Loan      24789
Personal Loan     4146
Contact Method    0
Campaign Calls    0
Previous Contact Days 0
Subscription Status 0
dtype: int64
```


Check the target column for distribution

Insights

There is a very high class imbalance, hence we need to consider SMOTE or UnderSampling when training our model.

```
In [11]: df['Subscription Status'].value_counts().plot(kind='bar')
df['Subscription Status'].value_counts()/df.shape[0]
```

```
Out[11]: Subscription Status
no      0.887346
yes    0.112654
Name: count, dtype: float64
```



(2a) Column by Column Data Understanding & Insights

Data Understanding & Insights for Client ID

From the description of Client ID, we can see that it is a unique identifier for each client. This identifier column does not contain any missing values (41188 non-null values), which is equal to the total number of records in the dataset (41188).

```
In [12]: display(df["Client ID"].nunique())
display(df["Client ID"].isna().sum())
display(df["Client ID"].describe())
```

```
41188
np.int64(0)
count    41188.00000
mean     20594.50000
std      11890.09578
min      1.00000
25%     10297.75000
50%     20594.50000
75%     30891.25000
max     41188.00000
Name: Client ID, dtype: float64
```

Data Understanding & Insights for Age

The "Age" feature in this dataset is stored as text, so it appears as an 'object' type.

There are 77 unique ages

The age "150 years" happens to appear in this dataset 4197 times, which is extremely unrealistic. The oldest person alive right now is 116 years old, which further solidifies the fact that there cannot be that many people who are significantly older.

Hence, we will treat this value as an outlier

The next oldest age after ""150 years" is "95 years", which is much more realistic.

```
In [13]: display(df["Age"].dtype)
display(df["Age"].value_counts())
```

```
dtype('O')
Age
150 years    4197
31 years     1747
32 years     1646
33 years     1643
36 years     1606
...
92 years      4
89 years      2
91 years      2
98 years      2
95 years      1
Name: count, Length: 77, dtype: int64
```

```
In [14]: print("The number of unique ages is",df["Age"].nunique())
```

The number of unique ages is 77

```
In [15]: display(df["Age"].value_counts().sort_index(ascending=False).head(10))
display(df["Age"].value_counts().sort_index(ascending=False).tail(10))
```

```
Age
98 years      2
95 years      1
92 years      4
91 years      2
89 years      2
88 years     19
86 years      8
85 years     15
84 years      7
83 years     16
Name: count, dtype: int64
Age
25 years     536
24 years     414
23 years     207
22 years     123
21 years      90
20 years      62
19 years      38
18 years      27
17 years       5
150 years    4197
Name: count, dtype: int64
```

Data Understanding & Insights for Occupation

This feature has a total of 12 different categories, with zero missing values.

As this class is a categorical feature, we can use one-hot encoding.

Later on during the bivariate analysis, we will be able to check how the different occupations can affect the subscription rates.


```
In [16]: print("The number of occupations in this dataset is",df["Occupation"].nunique())
df["Occupation"].value_counts()
```

The number of occupations in this dataset is 12

```
Out[16]: Occupation
admin.           10422
blue-collar     9254
technician      6743
services         3969
management      2924
retired          1720
entrepreneur    1456
self-employed    1421
housemaid        1060
unemployed       1014
student           875
unknown            330
Name: count, dtype: int64
```

```
In [17]: df['Occupation'].value_counts().plot(kind='bar')
df['Occupation'].info()
plt.show()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 41188 entries, 0 to 41187
Series name: Occupation
Non-Null Count   Dtype  
----- 
41188 non-null   object 
dtypes: object(1)
memory usage: 321.9+ KB
```



Data Understanding & Insights for Marital Status

This feature has a total of 4 different categories, with zero missing values.

Similarly to the "Occupation" column, this class is a categorical feature, so we can use one-hot encoding.

Later on during the bivariate analysis, we will be able to check how the different marital statuses can affect the subscription rates.


It is particularly useful, as marital status can definitely influence financial decisions due to differing priorities and interests.

```
In [18]: display(df["Marital Status"].value_counts())
display(df["Marital Status"].value_counts()/df.shape[0]*100)
```

```
Marital Status
married      24928
single       11568
divorced     4612
unknown        80
Name: count, dtype: int64
Marital Status
married      60.522482
single       28.085850
divorced     11.197436
unknown       0.194231
Name: count, dtype: float64
```

```
In [19]: df['Marital Status'].value_counts().plot(kind='pie', autopct='%1.1f%%')
df['Marital Status'].info()
plt.show()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 41188 entries, 0 to 41187
Series name: Marital Status
Non-Null Count Dtype
-----
41188 non-null object
dtypes: object(1)
memory usage: 321.9+ KB
```



Data Understanding & Insights for Education Level

This feature has a total of 8 different categories, with zero missing values.

This class is also a categorical feature, so we can use one-hot encoding.

Later on during the bivariate analysis, we will be able to check how the different education levels can affect the subscription rates.

Statistically, education generally correlates with income and financial literacy, which could affect the likelihood of individuals subscribing to the term deposit plan.

```
In [20]: print("The number of different Education Levels in this dataset is",df["Education Level"].value_counts())
```

The number of different Education Levels in this dataset is 8


```
Out[20]: Education Level
university.degree      12168
high.school            9515
basic.9y               6045
professional.course    5243
basic.4y                4176
basic.6y                2292
unknown                  1731
illiterate                 18
Name: count, dtype: int64
```

```
In [21]: df[df["Education Level"]=="illiterate"]["Subscription Status"].head()
```

```
Out[21]: 3824     no
5446     no
5742    yes
6342     no
8213     no
Name: Subscription Status, dtype: object
```

```
In [22]: df['Education Level'].value_counts().plot(kind='bar')
df['Education Level'].info()
plt.show()
```

```
<class 'pandas.core.series.Series'>
RangeIndex: 41188 entries, 0 to 41187
Series name: Education Level
Non-Null Count Dtype
-----
41188 non-null  object
dtypes: object(1)
memory usage: 321.9+ KB
```



Data Understanding & Insights for Credit Default

This feature has a total of 3 different categories, with zero missing values. However there is a decently large number of the "unknown" class

There is an extreme imbalance of data for this column, as seen below, where there are only three instances of "yes", which is 0.007% of the data.

We will need to consider how we use the "yes" class later on, due to its scarcity.

```
In [23]: display(df["Credit Default"].value_counts())
display(df["Credit Default"].value_counts()/df.shape[0]*100)
```

```
Credit Default
no        32588
unknown    8597
yes         3
Name: count, dtype: int64
Credit Default
no        79.120132
unknown   20.872584
yes       0.007284
Name: count, dtype: float64
```

Data Understanding & Insights for Housing Loan

This feature has a total of 3 different categories.

There is a large number of null values for this column (24789), which is ~60% of the data.

We will need to consider how we classify the null values later on, due to how much of the data it constitutes.

```
In [24]: display(df["Housing Loan"].isna().sum())
display(df["Housing Loan"].isna().sum()/df.shape[0]*100)
display(df["Housing Loan"].value_counts())
```

```
np.int64(24789)
np.float64(60.18500534136157)
Housing Loan
yes      8595
no       7411
unknown   393
Name: count, dtype: int64
```

Data Understanding & Insights for Personal Loan


This feature has a total of 3 different categories.

Similar to the previous column, Housing Loan, there are null values for this column (4146), which is ~10% of the data.

This feature will be useful as having personal loan(s) could indicate that the person already has financial commitments and this could reduce the likeliness of them subscribing to a new term deposit. We will need to consider how we classify the null values later on.

```
In [25]: display(df["Personal Loan"].isna().sum())
display(df["Personal Loan"].isna().sum()/df.shape[0]*100)
display(df["Personal Loan"].value_counts())
display(df["Personal Loan"].value_counts()/df.shape[0]*100)
display(df["Personal Loan"].value_counts().plot(kind='bar'))
```

```
np.int64(4146)
np.float64(10.066038652034573)
Personal Loan
no      30532
yes     5633
unknown  877
Name: count, dtype: int64
Personal Loan
no      74.128387
yes     13.676313
unknown  2.129261
Name: count, dtype: float64
<Axes: xlabel='Personal Loan'>
```



Data Understanding & Insights for Contact Method

This feature has a total of 4 different categories, with zero missing values.

This class is also a categorical feature, so we can use one-hot encoding.


However, it seems that the categories are inconsistent:

- Cell & cellular
- Telephone & telephone

Before encoding, we will standardise the data by combining these into fewer categories as they have the same meanings.

```
In [26]: df["Contact Method"].value_counts()  
df["Contact Method"].value_counts().plot(kind='bar')
```

```
Out[26]: <Axes: xlabel='Contact Method'>
```



Data Understanding & Insights for Campaign Calls

This feature has zero missing values.

By looking at the output from the describe function, we can make some interesting insights:

- The maximum value is 56 calls, meaning there are some outliers with heavy-contact
- The minimum value is -41 calls, which is definitely an error as it is impossible to have negative calls

Upon closer inspection, there are multiple values in this feature with a negative value. These invalid values need to be handled later by adjusting rows where the "Campaign Calls" is less than zero.

```
In [27]: display(df["Campaign Calls"].value_counts())
display(df["Campaign Calls"].describe())
```

```
Campaign Calls
1      15874
2      9446
3      4807
4      2405
-1     1768
...
56      1
-41     1
39      1
37      1
-25     1
Name: count, Length: 70, dtype: int64
count    41188.000000
mean      2.051374
std       3.171345
min      -41.000000
25%      1.000000
50%      2.000000
75%      3.000000
max      56.000000
Name: Campaign Calls, dtype: float64
```

```
In [28]: display(df["Campaign Calls"].value_counts().sort_index(ascending=False).head(10))
display(df["Campaign Calls"].value_counts().sort_index(ascending=False).tail(10))
```

```
Campaign Calls
56      1
43      2
42      2
40      2
39      1
37      1
35      3
34      3
33      4
32      3
Name: count, dtype: int64
Campaign Calls
-20     3
-21     2
-22     3
-23     5
-25     1
-28     2
-29     1
-32     1
-35     2
-41     1
Name: count, dtype: int64
```

Data Understanding & Insights for Previous Contact Days

This feature has zero missing values.

Majority of the values in this feature are "999"

The project guide has stated that "999" means there has been no previous contact

However leaving this as a numerical value of 999 will heavily skew and distort the actual statistics.

Hence, we will create a different feature to address this later.

```
In [29]: display(df["Previous Contact Days"].value_counts().sort_index(ascending=False).head())
display(df["Previous Contact Days"].value_counts().sort_index(ascending=False).tail())

Previous Contact Days
999    39673
27      1
26      1
25      1
22      3
21      2
20      1
19      3
18      7
17      8
Name: count, dtype: int64
Previous Contact Days
9      64
8      18
7      60
6     412
5      46
4     118
3     439
2      61
1      26
0      15
Name: count, dtype: int64
```

Further Insights for the distribution of Subscription Status

Looking further into this feature, we can see that the no to yes ratio nearly reaches a 9:1.

This shows that this dataset is heavily imbalanced and skewed.

When training our models during the pipeline, we definitely take some factors into consideration, such as:

- Proper sampling
- Appropriate metrics
- Assigning weights to the data

```
In [30]: df["Subscription Status"].value_counts()
percentage_wise=df["Subscription Status"].value_counts()/df.shape[0]*100
print(percentage_wise)
```

```
Subscription Status
no      88.734583
yes     11.265417
Name: count, dtype: float64
```

(3) Handling the null values / incorrect values / outliers

In order to convert the string inside the "Age" column into integer values, we must do the following:

```
In [31]: df["Age_numerical"]=(df["Age"].str.replace('years', '', regex=False)).astype('int')
```

For this next part, we addressed the outlier age of "150", by imputing the value using the median age. Firstly, we are commonly taught to drop the values if it is 5% of the total number of observations. However the total number of instances where the age is equal to 150 is ~10%, hence we decided on using imputation of values. Next we had to choose between median and mean imputation. Median imputation was chosen as the ages goes all the way up to 98 and median is more stable than mean, as it will not skew the cleaned dataset as much.

```
In [32]: df_150=df[df["Age"]=="150 years"]
print(df_150.head(50))
```

| | Client ID | Age | Occupation | Marital Status | Education Level | \ |
|-----|-----------|-----------|---------------|----------------|---------------------|---|
| 7 | 23758 | 150 years | admin. | divorced | university.degree | |
| 13 | 21872 | 150 years | admin. | married | university.degree | |
| 17 | 17689 | 150 years | technician | married | professional.course | |
| 51 | 38203 | 150 years | retired | married | high.school | |
| 59 | 29732 | 150 years | blue-collar | married | basic.4y | |
| 61 | 4501 | 150 years | admin. | single | university.degree | |
| 67 | 1632 | 150 years | admin. | married | high.school | |
| 73 | 35637 | 150 years | services | single | high.school | |
| 77 | 24445 | 150 years | admin. | single | high.school | |
| 82 | 4244 | 150 years | blue-collar | married | basic.4y | |
| 83 | 13927 | 150 years | services | single | high.school | |
| 89 | 8086 | 150 years | blue-collar | married | basic.4y | |
| 94 | 38325 | 150 years | services | married | basic.9y | |
| 106 | 2001 | 150 years | self-employed | married | university.degree | |
| 120 | 11097 | 150 years | technician | single | professional.course | |
| 125 | 22364 | 150 years | blue-collar | married | basic.4y | |
| 127 | 13029 | 150 years | admin. | single | professional.course | |
| 129 | 10099 | 150 years | services | married | high.school | |
| 135 | 15045 | 150 years | technician | divorced | high.school | |
| 143 | 35754 | 150 years | admin. | divorced | high.school | |
| 144 | 9591 | 150 years | technician | single | basic.9y | |
| 147 | 22847 | 150 years | technician | married | university.degree | |
| 150 | 37978 | 150 years | technician | married | professional.course | |
| 164 | 13416 | 150 years | technician | married | university.degree | |
| 166 | 33667 | 150 years | self-employed | single | professional.course | |
| 183 | 20191 | 150 years | self-employed | married | university.degree | |
| 189 | 40164 | 150 years | retired | divorced | basic.4y | |
| 199 | 9255 | 150 years | technician | married | basic.9y | |
| 212 | 24204 | 150 years | unemployed | married | university.degree | |
| 219 | 29527 | 150 years | blue-collar | married | basic.9y | |
| 223 | 41101 | 150 years | admin. | married | university.degree | |
| 226 | 2005 | 150 years | entrepreneur | married | high.school | |
| 246 | 37873 | 150 years | technician | single | professional.course | |
| 270 | 9507 | 150 years | management | married | basic.4y | |
| 279 | 14929 | 150 years | blue-collar | married | basic.4y | |
| 294 | 26167 | 150 years | management | single | high.school | |
| 304 | 1431 | 150 years | blue-collar | married | basic.6y | |
| 305 | 13207 | 150 years | services | single | high.school | |
| 306 | 21013 | 150 years | technician | married | professional.course | |
| 315 | 11844 | 150 years | admin. | divorced | high.school | |
| 317 | 24855 | 150 years | admin. | married | university.degree | |
| 319 | 38125 | 150 years | management | married | university.degree | |
| 322 | 37512 | 150 years | blue-collar | married | basic.9y | |
| 344 | 28487 | 150 years | self-employed | married | university.degree | |
| 373 | 20413 | 150 years | admin. | married | university.degree | |
| 378 | 12958 | 150 years | blue-collar | divorced | basic.9y | |
| 382 | 10931 | 150 years | admin. | married | university.degree | |
| 392 | 28752 | 150 years | blue-collar | married | basic.4y | |
| 395 | 24511 | 150 years | housemaid | married | university.degree | |
| 402 | 5750 | 150 years | services | single | basic.9y | |

| | Credit | Default | Housing | Loan | Personal | Loan | Contact | Method | Campaign | Calls | \ |
|----|---------|---------|---------|------|----------|------|-----------|--------|----------|-------|---|
| 7 | unknown | | yes | | None | | Cell | | 7 | | |
| 13 | no | | no | | no | | Cell | | 3 | | |
| 17 | no | | None | | yes | | cellular | | -11 | | |
| 51 | no | | None | | no | | cellular | | 1 | | |
| 59 | unknown | | None | | no | | cellular | | 1 | | |
| 61 | no | | None | | no | | telephone | | 1 | | |
| 67 | no | | None | | no | | Telephone | | 2 | | |

| | | | | | |
|-----|---------|------|---------|-----------|----|
| 73 | no | None | no | Cell | 6 |
| 77 | no | None | no | cellular | 1 |
| 82 | no | yes | yes | Telephone | 2 |
| 83 | no | None | None | Cell | 2 |
| 89 | unknown | None | no | Telephone | 2 |
| 94 | no | None | no | Telephone | 1 |
| 106 | no | None | no | telephone | 5 |
| 120 | unknown | None | yes | Telephone | 1 |
| 125 | no | None | no | Cell | 1 |
| 127 | no | yes | yes | Cell | 3 |
| 129 | unknown | no | no | telephone | 1 |
| 135 | no | no | no | Cell | 1 |
| 143 | no | None | no | Telephone | 2 |
| 144 | no | None | no | Telephone | 3 |
| 147 | no | None | no | cellular | 2 |
| 150 | no | None | no | cellular | 1 |
| 164 | no | None | unknown | Telephone | 1 |
| 166 | no | None | no | cellular | 2 |
| 183 | no | None | None | cellular | 2 |
| 189 | no | None | no | Cell | 2 |
| 199 | no | no | no | Telephone | -2 |
| 212 | no | None | unknown | Cell | 1 |
| 219 | unknown | None | yes | Cell | 5 |
| 223 | no | None | no | cellular | 1 |
| 226 | no | yes | no | Telephone | 2 |
| 246 | no | yes | no | Cell | 1 |
| 270 | no | no | None | telephone | 2 |
| 279 | unknown | yes | no | Cell | 4 |
| 294 | no | yes | yes | Cell | 2 |
| 304 | unknown | None | no | Telephone | 2 |
| 305 | no | None | no | cellular | -1 |
| 306 | no | no | yes | Cell | 2 |
| 315 | no | no | no | telephone | 3 |
| 317 | no | no | no | cellular | 1 |
| 319 | no | yes | no | cellular | 1 |
| 322 | no | None | no | cellular | 1 |
| 344 | no | None | no | Cell | 2 |
| 373 | no | None | yes | Cell | 1 |
| 378 | no | None | no | cellular | 1 |
| 382 | no | yes | no | Telephone | 5 |
| 392 | no | None | no | cellular | 1 |
| 395 | no | None | no | cellular | 1 |
| 402 | no | None | no | telephone | -2 |

| | Previous_Contact | Days | Subscription_Status | Age_numerical |
|-----|------------------|------|---------------------|---------------|
| 7 | | 999 | no | 150 |
| 13 | | 999 | yes | 150 |
| 17 | | 999 | no | 150 |
| 51 | | 999 | no | 150 |
| 59 | | 999 | no | 150 |
| 61 | | 999 | no | 150 |
| 67 | | 999 | no | 150 |
| 73 | | 999 | no | 150 |
| 77 | | 999 | no | 150 |
| 82 | | 999 | no | 150 |
| 83 | | 999 | no | 150 |
| 89 | | 999 | no | 150 |
| 94 | | 999 | yes | 150 |
| 106 | | 999 | no | 150 |
| 120 | | 999 | no | 150 |

| | | | |
|-----|-----|-----|-----|
| 125 | 999 | no | 150 |
| 127 | 999 | no | 150 |
| 129 | 999 | no | 150 |
| 135 | 999 | no | 150 |
| 143 | 999 | no | 150 |
| 144 | 999 | no | 150 |
| 147 | 999 | no | 150 |
| 150 | 999 | no | 150 |
| 164 | 999 | no | 150 |
| 166 | 999 | yes | 150 |
| 183 | 999 | no | 150 |
| 189 | 6 | yes | 150 |
| 199 | 999 | no | 150 |
| 212 | 999 | no | 150 |
| 219 | 999 | no | 150 |
| 223 | 12 | yes | 150 |
| 226 | 999 | no | 150 |
| 246 | 999 | no | 150 |
| 270 | 999 | no | 150 |
| 279 | 999 | no | 150 |
| 294 | 6 | no | 150 |
| 304 | 999 | no | 150 |
| 305 | 999 | no | 150 |
| 306 | 999 | no | 150 |
| 315 | 999 | no | 150 |
| 317 | 999 | no | 150 |
| 319 | 999 | no | 150 |
| 322 | 999 | no | 150 |
| 344 | 999 | yes | 150 |
| 373 | 999 | no | 150 |
| 378 | 999 | no | 150 |
| 382 | 999 | no | 150 |
| 392 | 999 | no | 150 |
| 395 | 999 | no | 150 |
| 402 | 999 | no | 150 |

```
In [33]: df["Age_Invalid"] = (df["Age_numerical"] == 150).astype('int')
valid_median_age=df.loc[df["Age_numerical"]!=150,"Age_numerical"].median()
df.loc[df["Age_numerical"]==150,"Age_numerical"]= valid_median_age
df["Age_numerical"].describe()
```

```
Out[33]: count    41188.000000
mean      39.818928
std       9.909692
min      17.000000
25%      33.000000
50%      38.000000
75%      46.000000
max      98.000000
Name: Age_numerical, dtype: float64
```

Handling of unknown values in the "Occupation" Column

When inspecting the "Housing Loan" column, we can see that there are many different categories.

We have decided to remove the unknown values here alone as it constitutes to a very small minority of the data and we cannot combine it with another feature as there is no correlation.

Unlike unknown, we can keep the other categories, as we can combine them into larger groups later on for analysis.

The difference between unknown and the rest, is that unknown carries 0 value by itself

```
In [34]: df["Occupation"].value_counts()
```

```
Out[34]: Occupation
admin.          10422
blue-collar     9254
technician      6743
services         3969
management      2924
retired          1720
entrepreneur    1456
self-employed    1421
housemaid        1060
unemployed       1014
student           875
unknown            330
Name: count, dtype: int64
```

```
In [35]: df=df[df["Occupation"]!="unknown"]
df["Occupation"].value_counts()
```

```
Out[35]: Occupation
admin.          10422
blue-collar     9254
technician      6743
services         3969
management      2924
retired          1720
entrepreneur    1456
self-employed    1421
housemaid        1060
unemployed       1014
student           875
Name: count, dtype: int64
```

Handling of unknown values in the "Marital Status" Column

When inspecting the "Housing Loan" column, we can see that there are 4 different categories, "married", "single", "unknown" and "divorced".

We have decided to remove the unknown values here alone as it constitutes to a very small minority of the data.

```
In [36]: df=df[df["Marital Status"]!="unknown"]
df["Marital Status"].value_counts()
```

```
Out[36]: Marital Status
married          24694
single           11494
divorced          4599
Name: count, dtype: int64
```

Handling of unknown values in the "Education Level" Column

When inspecting the "Education Level" column, we can see that there are multiple different categories: university degree, high school, basic (4 years), basic (6 years), basic (9 years), professional course, unknown and illiterate.

A major thing we took note of was the uneven distribution of clients in the different categories for this column, and further inspection into the unknown and illiterate categories show that they total to less than 5% of the total data.

Although it is a good choice to have the "unknown" and "illiterate" columns, they face an extreme class imbalance. There are only a total of 1749 instances of these columns having data, and this is out of a total of around 40k. This means that it only constitutes to around 4.25%, which is quite low. Hence, these 2 categories would not affect training or evaluation of predictive models as it does not provide much meaning or value to prediction. These categories also cannot be merged with any other column, as they are each unique and have completely different meanings. Hence, we decided to exclude these 2 categories from the final cleaned dataset. This will be applied to the training of the selected models to prevent overfitting, as keeping these features would cause instability besides from not contributing any predictive value.

```
In [37]: display(df["Education Level"].value_counts())
display(df["Education Level"].unique())
```

```
Education Level
university.degree    12096
high.school          9464
basic.9y             6006
professional.course  5225
basic.4y              4118
basic.6y              2264
unknown               1596
illiterate                18
Name: count, dtype: int64
array(['high.school', 'basic.9y', 'professional.course',
       'university.degree', 'basic.4y', 'basic.6y', 'unknown',
       'illiterate'], dtype=object)
```

```
In [38]: df=df[df["Education Level"]!="illiterate"]
df=df[df["Education Level"]!="unknown"]
```

```
In [39]: df["Education Level"].value_counts()
```

```
Out[39]: Education Level
university.degree    12096
high.school          9464
basic.9y             6006
professional.course  5225
basic.4y              4118
basic.6y              2264
Name: count, dtype: int64
```

Handling of the "Credit Default" Features

When inspecting the "Credit Default" column, we can see that there are 3 different categories, "yes", "no" and "unknown"

No credit default can indicate that the client has financial stability and is more

responsible with their money, which can affect their decision on whether they should or should not subscribe. The "unknown" class in this data indicate that the client has not have borrowed money from this bank before or the record could be incomplete. "Yes" in the context of credit default means that they did not manage to make the payments on time, and could signify that the person is tight on money. Although it is a good choice to have the "Yes" column, it faces an extreme class imbalance. There are only 3 instances of this being positive, and this is out of a total of 41188. This means that it only constitutes to around 0.00728%, which is extremely low. Hence, the "Yes" category would not affect training or evaluation of predictive models as it does not provide any meaning or value towards predictions. This category also cannot be merged with "unknown", as they have completely different meanings. Hence, we decided to exclude this category from the final cleaned dataset. This will be applied to the training of the selected models to prevent overfitting, as keeping this feature would cause instability besides from not contributing any predictive value.

```
In [40]: df["Credit Default"].value_counts()
```

```
Out[40]: Credit Default
no           31213
unknown      7957
yes          3
Name: count, dtype: int64
```

```
In [41]: df=df[df["Credit Default"]!="yes"]
df["Credit Default"].value_counts()
```


```
Out[41]: Credit Default
no           31213
unknown      7957
Name: count, dtype: int64
```

Before moving on to the next features, we wanted to visualize and confirm that the next two rows we would be working on do indeed have null values.

```
In [42]: import missingno as msno
```

```
In [43]: msno.bar(df)
```

```
Out[43]: <Axes: >
```



```
In [44]: pd.options.display.max_rows = 999
nulls = df.isna().sum().sort_values(ascending=False)
nulls=nulls>0
```

```
Out[44]: Housing Loan      23566
          Personal Loan     3940
          dtype: int64
```

```
In [45]: df.shape
```

```
Out[45]: (39170, 14)
```

Cleaning of the "Housing Loan" Feature

When inspecting the "Housing Loan" column, we can see that there are 3 different categories, "yes", "no", "unknown" and the null values appear as as "None". However, "None" and "no" do not mean the same thing in this case. "no" most likely refers to the client not having any housing loans, while "None" means that the bank has never recorded down the status of the client's housing loans, hence we are not able to verify whether they have or do not have any housing loans.

We have decided to remove the unknown values here alone as it constitutes to a very small minority of the data.

```
In [46]: df[df["Housing Loan"].isna()].head(15)
```

Out[46]:

| | Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Personal Loan |
|----|-----------|-----------|--------------|----------------|---------------------|----------------|--------------|---------------|
| 4 | 14021 | 27 years | housemaid | married | high.school | no | None | no |
| 5 | 17202 | 58 years | retired | married | professional.course | no | None | yes |
| 8 | 10822 | 24 years | entrepreneur | married | university.degree | no | None | None |
| 10 | 32312 | 34 years | blue-collar | married | basic.9y | no | None | no |
| 12 | 13596 | 43 years | blue-collar | single | basic.9y | no | None | no |
| 14 | 16736 | 58 years | blue-collar | married | basic.4y | no | None | yes |
| 17 | 17689 | 150 years | technician | married | professional.course | no | None | yes |
| 18 | 15508 | 37 years | entrepreneur | single | professional.course | no | None | no |
| 20 | 17268 | 59 years | retired | married | basic.9y | unknown | None | no |
| 21 | 1671 | 49 years | unemployed | single | professional.course | unknown | None | no |
| 23 | 15756 | 30 years | admin. | single | high.school | no | None | yes |
| 25 | 36916 | 45 years | admin. | married | basic.9y | no | None | no |
| 27 | 22970 | 31 years | admin. | married | university.degree | no | None | no |
| 28 | 8926 | 41 years | blue-collar | married | basic.4y | unknown | None | no |
| 30 | 36969 | 47 years | housemaid | married | basic.6y | unknown | None | yes |

In [47]: `df[df["Housing Loan"]=="no"].head(15)`

Out[47]:

| | Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Person Lo |
|-----|-----------|-----------|--------------|----------------|---------------------|----------------|--------------|-----------|
| 0 | 32885 | 57 years | technician | married | high.school | no | no | y |
| 2 | 32207 | 33 years | blue-collar | married | basic.9y | no | no | |
| 3 | 9404 | 36 years | admin. | married | high.school | no | no | |
| 13 | 21872 | 150 years | admin. | married | university.degree | no | no | |
| 24 | 20771 | 51 years | management | married | high.school | unknown | no | |
| 31 | 36563 | 69 years | entrepreneur | married | high.school | no | no | |
| 38 | 8726 | 30 years | blue-collar | married | basic.6y | no | no | |
| 39 | 25810 | 30 years | management | married | university.degree | no | no | |
| 64 | 40577 | 33 years | technician | single | professional.course | no | no | |
| 70 | 18429 | 47 years | admin. | married | high.school | unknown | no | y |
| 92 | 25159 | 40 years | admin. | divorced | basic.9y | no | no | |
| 98 | 533 | 41 years | blue-collar | married | basic.6y | no | no | |
| 108 | 17076 | 42 years | admin. | single | high.school | no | no | y |
| 110 | 29676 | 33 years | blue-collar | married | high.school | no | no | |
| 123 | 6223 | 36 years | services | married | high.school | unknown | no | |

In [48]: `df[df["Housing Loan"]=="unknown"].head(15)`

Out[48]:

| Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Personal Lo |
|-----------|-------|------------|----------------|-----------------|---------------------|--------------|-------------|
| 169 | 36548 | 50 years | admin. | single | basic.9y | no | unknown |
| 211 | 14241 | 27 years | blue-collar | single | basic.9y | no | unknown |
| 244 | 14423 | 32 years | admin. | single | university.degree | no | unknown |
| 401 | 31694 | 38 years | technician | married | university.degree | no | unknown |
| 404 | 6661 | 54 years | admin. | married | high.school | no | unknown |
| 468 | 9332 | 40 years | admin. | married | university.degree | unknown | unknown |
| 804 | 15179 | 29 years | unemployed | married | university.degree | no | unknown |
| 860 | 26664 | 34 years | unemployed | married | high.school | no | unknown |
| 906 | 19859 | 49 years | technician | married | professional.course | unknown | unknown |
| 952 | 20069 | 43 years | admin. | married | university.degree | no | unknown |
| 962 | 26267 | 32 years | admin. | married | professional.course | no | unknown |
| 1206 | 1034 | 40 years | blue-collar | married | basic.4y | unknown | unknown |
| 1242 | 6003 | 49 years | blue-collar | married | basic.4y | no | unknown |
| 1311 | 29585 | 35 years | admin. | single | university.degree | no | unknown |
| 1328 | 13421 | 150 years | technician | married | university.degree | no | unknown |

In [49]: `display(df["Housing Loan"].value_counts())`

```
Housing Loan
yes      8193
no       7035
unknown    376
Name: count, dtype: int64
```

In [50]: `df=df[df["Housing Loan"]!="unknown"]`

```
df[ "Personal Loan"].value_counts()
df[ "Housing Loan"]=df[ "Housing Loan"].fillna("no_info")
```

```
display(df["Housing Loan"].isna().value_counts())
display(df["Housing Loan"].value_counts())
```

```
Housing Loan
False    38794
Name: count, dtype: int64
Housing Loan
no_info    23566
yes        8193
no         7035
Name: count, dtype: int64
```

Similar to the "Housing Loan" column, there are 3 different categories, "yes", "no", "unknown" and null values as "None".

However, "None" and "no" do not mean the same thing in this case. "no" most likely refers to the client not having any personal loans, while "None" means that the bank has never recorded down the status of the client's personal loans, hence we are not able to verify whether they have or do not have any personal loans.

We have decided to remove the unknown values as it constitutes to a very small minority of the data.

```
In [51]: display(df["Personal Loan"].value_counts())
```

```
Personal Loan
no        29020
yes       5370
unknown    509
Name: count, dtype: int64
```

```
In [52]: display(df[df["Personal Loan"].isna()].head(5))
display(df[df["Personal Loan"]=="no"].head(5))
df[df["Personal Loan"]=="unknown"].head(5)
```

| | Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Personal Loan |
|----|-----------|-----------|--------------|----------------|-------------------|----------------|--------------|---------------|
| 7 | 23758 | 150 years | admin. | divorced | university.degree | unknown | yes | None |
| 8 | 10822 | 24 years | entrepreneur | married | university.degree | no | no_info | None |
| 58 | 39197 | 41 years | admin. | divorced | university.degree | no | yes | None |
| 63 | 3420 | 45 years | blue-collar | married | basic.4y | no | no_info | None |
| 80 | 33318 | 41 years | management | married | high.school | unknown | yes | None |



| Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Personal Loan | Contact Method |
|-----------|----------|-------------|----------------|-----------------|----------------|--------------|---------------|----------------|
| 2 32207 | 33 years | blue-collar | married | basic.9y | no | no | no | cellular |
| 3 9404 | 36 years | admin. | married | high.school | no | no | no | Telephone |
| 4 14021 | 27 years | housemaid | married | high.school | no | no_info | no | Cellular |
| 6 880 | 48 years | services | married | high.school | unknown | yes | no | Telephone |
| 10 32312 | 34 years | blue-collar | married | basic.9y | no | no_info | no | Cellular |



Out[52]:

| Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Personal Loan |
|-----------|-----------|------------|----------------|---------------------|----------------|--------------|---------------|
| 164 13416 | 150 years | technician | married | university.degree | no | no_info | unknown |
| 188 8815 | 30 years | technician | single | professional.course | no | no_info | unknown |
| 212 24204 | 150 years | unemployed | married | university.degree | no | no_info | unknown |
| 225 26419 | 38 years | admin. | married | high.school | no | no_info | unknown |
| 239 8202 | 38 years | student | single | university.degree | no | no_info | unknown |



In [53]:

```
df=df[df["Personal Loan"]!="unknown"]
display(df["Personal Loan"].isna().sum())
df["Personal Loan"]=df["Personal Loan"].fillna("no_info")
display(df["Personal Loan"].isna().value_counts())
display(df["Personal Loan"].value_counts())
```

```
np.int64(3895)
Personal Loan
False      38285
Name: count, dtype: int64
Personal Loan
no          29020
yes         5370
no_info     3895
Name: count, dtype: int64
```

Cleaning of the "Contact Method" Feature

When inspecting the "Contact Method" column, we can see that there are 4 different categories, "Cell", "cellular", "Telephone" and "telephone".

To clean this, we will convert them into similar, lowercase terms:

"Telephone" & "telephone" --> "telephone"

"Cell" & "cellular" --> "cellular"

```
In [54]: df["Contact Method"].value_counts()
```

```
Out[54]: Contact Method
```

| | |
|-----------|-------|
| Cell | 12257 |
| cellular | 12201 |
| Telephone | 6933 |
| telephone | 6894 |

Name: count, dtype: int64

```
In [55]: df["Contact Method"] = df["Contact Method"].replace({"Cell": "cellular", "Telephon  
df["Contact Method"].value_counts()
```

```
Out[55]: Contact Method
```

| | |
|-----------|-------|
| cellular | 24458 |
| telephone | 13827 |

Name: count, dtype: int64

Cleaning of the "Campaign Calls" Feature

When inspecting the "Campaign" column, we can see that there are many values that are negative.

To clean this, we must first check the number of negative values, before we make a decision.

```
In [56]: df[df["Campaign Calls"] < 0].shape[0]
```

```
Out[56]: 3847
```

```
In [57]: df["Campaign Calls"].value_counts().sort_index(ascending=True).head(50)
```

```
Out[57]: Campaign Calls
```

```
-41      1
-35      2
-32      1
-29      1
-28      2
-25      1
-23      5
-22      3
-21      2
-20      3
-19      1
-18      3
-17      5
-16      5
-15      4
-14     10
-13      8
-12     15
-11     19
-10     15
-9      27
-8      34
-7      59
-6      82
-5    144
-4    231
-3    486
-2   1051
-1   1627
1    14759
2    8781
3    4473
4    2238
5    1360
6    830
7    525
8    339
9    232
10   193
11   147
12   100
13    68
14    54
15    43
16    40
17    49
18    28
19    22
20    27
21    17
```

```
Name: count, dtype: int64
```

```
In [58]: df["Campaign Calls"].value_counts().sort_index(ascending=False).head(50)
```

```
Out[58]: Campaign Calls
```

```
43      2  
42      2  
40      2  
39      1  
37      1  
35      3  
34      3  
33      4  
32      3  
31      7  
30      7  
29      8  
28      6  
27     11  
26      7  
25      6  
24     15  
23     11  
22     14  
21     17  
20     27  
19     22  
18     28  
17     49  
16     40  
15     43  
14     54  
13     68  
12    100  
11    147  
10    193  
9     232  
8     339  
7     525  
6     830  
5    1360  
4    2238  
3    4473  
2    8781  
1   14759  
-1   1627  
-2   1051  
-3   486  
-4   231  
-5   144  
-6    82  
-7    59  
-8    34  
-9    27  
-10   15
```

```
Name: count, dtype: int64
```

```
In [59]: df[df["Campaign Calls"] < 0].sort_index(ascending=True).head(50)
```

Out[59]:

| | Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Person Lo |
|-----|-----------|-----------|---------------|----------------|---------------------|----------------|--------------|-----------|
| 17 | 17689 | 150 years | technician | married | professional.course | no | no_info | y |
| 23 | 15756 | 30 years | admin. | single | high.school | no | no_info | y |
| 36 | 23686 | 33 years | admin. | single | university.degree | no | yes | |
| 48 | 23462 | 44 years | technician | married | professional.course | no | yes | |
| 52 | 2465 | 36 years | blue-collar | married | basic.9y | unknown | no_info | |
| 56 | 24130 | 52 years | technician | divorced | university.degree | no | no_info | |
| 66 | 38247 | 77 years | retired | married | university.degree | no | no_info | y |
| 69 | 37778 | 57 years | retired | married | university.degree | no | yes | |
| 70 | 18429 | 47 years | admin. | married | high.school | unknown | no | y |
| 72 | 32715 | 51 years | entrepreneur | married | professional.course | unknown | no_info | |
| 84 | 3271 | 46 years | blue-collar | single | basic.4y | no | no_info | |
| 96 | 28108 | 44 years | management | divorced | university.degree | no | yes | |
| 105 | 16549 | 49 years | self-employed | divorced | university.degree | no | no_info | y |
| 108 | 17076 | 42 years | admin. | single | high.school | no | no | y |
| 113 | 7946 | 54 years | admin. | married | high.school | unknown | yes | |
| 139 | 22233 | 45 years | technician | single | university.degree | no | no | |
| 146 | 34534 | 26 years | blue-collar | single | basic.9y | no | no | |
| 173 | 6113 | 53 years | self-employed | married | university.degree | no | no_info | y |
| 177 | 19566 | 41 years | unemployed | divorced | university.degree | unknown | no_info | |
| 197 | 26891 | 43 years | technician | married | professional.course | no | no | |

| Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Person Loan |
|-----------|-------|------------|----------------|-----------------|---------------------|--------------|-------------|
| 199 | 9255 | 150 years | technician | married | basic.9y | no | no |
| 250 | 35029 | 32 years | admin. | divorced | high.school | no | no_info |
| 259 | 33470 | 35 years | admin. | single | basic.9y | no | no_info |
| 260 | 34182 | 27 years | self-employed | married | basic.9y | no | yes |
| 273 | 37483 | 62 years | retired | married | university.degree | unknown | no_info |
| 275 | 22783 | 50 years | housemaid | married | basic.4y | no | no |
| 278 | 21100 | 38 years | technician | divorced | university.degree | unknown | no_info |
| 280 | 15639 | 30 years | management | single | university.degree | no | no_info |
| 293 | 23085 | 55 years | technician | divorced | university.degree | no | no_info |
| 303 | 11811 | 40 years | housemaid | married | basic.4y | unknown | no_info |
| 305 | 13207 | 150 years | services | single | high.school | no | no_info |
| 312 | 30880 | 39 years | blue-collar | married | basic.9y | no | no_info |
| 321 | 12514 | 47 years | admin. | married | basic.9y | no | yes |
| 325 | 17378 | 44 years | blue-collar | single | high.school | unknown | no_info |
| 330 | 15918 | 57 years | admin. | married | university.degree | no | yes |
| 348 | 1078 | 29 years | admin. | single | university.degree | no | no_info |
| 349 | 6189 | 55 years | admin. | married | high.school | no | no_info |
| 359 | 37701 | 32 years | admin. | married | university.degree | no | no_info |
| 368 | 29911 | 41 years | blue-collar | married | professional.course | no | no_info |
| 391 | 2993 | 26 years | unemployed | single | basic.9y | no | no_info |

| | Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Person Lo |
|-----|-----------|-----------|------------|----------------|---------------------|----------------|--------------|-----------|
| 397 | 9286 | 48 years | management | married | basic.9y | no | yes | |
| 402 | 5750 | 150 years | services | single | basic.9y | no | no_info | |
| 403 | 41162 | 33 years | admin. | married | university.degree | no | yes | |
| 414 | 31991 | 38 years | technician | married | university.degree | no | no_info | |
| 419 | 586 | 150 years | admin. | single | university.degree | no | no_info | y |
| 430 | 8553 | 58 years | services | married | basic.4y | no | no_info | |
| 443 | 37055 | 45 years | admin. | married | university.degree | no | no_info | |
| 448 | 21742 | 30 years | technician | single | professional.course | no | yes | y |
| 449 | 5458 | 31 years | services | single | basic.6y | no | no_info | |
| 464 | 25542 | 38 years | management | married | professional.course | unknown | no_info | |

After looking at the values, ~10% of the data is less than 0, so it is not wise to drop these rows.

Instead, we decided to impute the values, similar to what we did for the cleaning of the "Age" column.

We decided to use the absolute value method, where we convert the negative values into positives, as there are no set boundaries for this data column, and change -41 to 41 would still be less than the maximum value of the data.

Also, the distribution of the data does follow the general trend, where there is an inversely proportional relationship between the number of calls and number of clients.

```
In [60]: df["Campaign Calls"] = df["Campaign Calls"].abs()
display(df["Campaign Calls"].describe())
```

```
count    38285.000000
mean      2.566984
std       2.766944
min       1.000000
25%      1.000000
50%      2.000000
75%      3.000000
max     43.000000
Name: Campaign Calls, dtype: float64
```

```
In [61]: Q1 = df["Campaign Calls"].quantile(0.25)
Q3 = df["Campaign Calls"].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
print(f"IQR Bounds for 'Campaign Calls': Lower Bound = {lower_bound}, Upper Bound = {upper_bound}")
```

```
IQR Bounds for 'Campaign Calls': Lower Bound = -2.0, Upper Bound = 6.0
```

Cleaning of the "Previous Contact Days" Feature

As seen in the project guide, values like "999" are placeholders for "No previous contact".

Hence, we cannot treat this as a numeric value, or it will skew the other data.

We decided to replace 999 with NaN and create a new column with binary features to show whether there has been contact or not, using "1" and "0".

This helps to remove any instances which could heavily imbalance our data, and still preserve the knowledge of knowing which clients have had contact or no previous contact.

Keeping this information is extremely important for our project, as it can influence a person's decision.

```
In [62]: df["no_previous_contact"] = (df["Previous Contact Days"] == 999)
df["no_previous_contact"] = df["no_previous_contact"].astype(int)
df["Previous Contact Days_clean"] = df["Previous Contact Days"].replace(999, np.nan)
```

```
In [63]: display(df["Previous Contact Days"].head())
display(df["no_previous_contact"].head())
display(df["Previous Contact Days_clean"].head())
```

```
0    999
2    999
3    999
4    999
5    999
Name: Previous Contact Days, dtype: int64
0     1
2     1
3     1
4     1
5     1
Name: no_previous_contact, dtype: int64
0    NaN
2    NaN
3    NaN
4    NaN
5    NaN
Name: Previous Contact Days_clean, dtype: float64
```

```
In [64]: df.head()
```

Out[64]:


| | Client ID | Age | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Personal Loan |
|---|-----------|----------|-------------|----------------|---------------------|----------------|--------------|---------------|
| 0 | 32885 | 57 years | technician | married | high.school | no | no | yes |
| 2 | 32207 | 33 years | blue-collar | married | basic.9y | no | no | no |
| 3 | 9404 | 36 years | admin. | married | high.school | no | no | no |
| 4 | 14021 | 27 years | housemaid | married | high.school | no | no_info | no |
| 5 | 17202 | 58 years | retired | married | professional.course | no | no_info | yes |



In [65]: `df=df.reset_index(drop=True)`

In [66]: `num_cols = ["Age_numerical", "Campaign Calls", "Previous Contact Days_clean"]`

```
plt.figure(figsize=(30,15))
for i, col in enumerate(num_cols, 1):
    plt.subplot(1, 3, i)
    sns.histplot(df[col], kde=True)
    plt.title(f'Distribution of {col}')
plt.tight_layout()
plt.show()
```



(4) Bivariate Analysis

Our main target variable for this EDA is "Subscription Status". However It is currently categorical (yes/no)

In order to conduct this bivariate analysis, we first have to encode it into numeric characters in order for this to work.

The following changes will be made:

```
"1"-> subscribed ("yes")
"0"-> not subscribed ("no")
```

```
In [67]: df["Subscription_Status_Code"] = df["Subscription Status"].map({"yes":1, "no":0})
df["Subscription_Status_Code"].value_counts()
```


```
Out[67]: Subscription_Status_Code
0    34025
1    4260
Name: count, dtype: int64
```

```
In [68]: df_cleaned = df.drop(columns=["Age", "Previous Contact Days"])
```

```
In [69]: df_cleaned.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38285 entries, 0 to 38284
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Client ID        38285 non-null   int64  
 1   Occupation       38285 non-null   object  
 2   Marital Status   38285 non-null   object  
 3   Education Level  38285 non-null   object  
 4   Credit Default   38285 non-null   object  
 5   Housing Loan     38285 non-null   object  
 6   Personal Loan    38285 non-null   object  
 7   Contact Method   38285 non-null   object  
 8   Campaign Calls   38285 non-null   int64  
 9   Subscription Status 38285 non-null   object  
 10  Age_numerical   38285 non-null   int64  
 11  Age_Invalid     38285 non-null   int64  
 12  no_previous_contact 38285 non-null   int64  
 13  Previous Contact Days_clean 1366 non-null   float64 
 14  Subscription_Status_Code 38285 non-null   int64  
dtypes: float64(1), int64(6), object(8)
memory usage: 4.4+ MB
```

```
In [70]: plt.figure(figsize=(10,6))
num_cols= ["Campaign Calls", "Age_numerical", "Age_Invalid", "no_previous_contact"]
df_numerical_relationship=df_cleaned[num_cols]
sns.heatmap(df_numerical_relationship.corr(numeric_only=True), annot=True, fmt="."
plt.title("Correlation Matrix for Numerical Features")
plt.show()
```






```
In [71]: sns.pairplot(df_cleaned)
```

```
Out[71]: <seaborn.axisgrid.PairGrid at 0x140fbb8a3c0>
```




```
In [72]: continuous_features = ["Age_numerical", "Campaign Calls", "Previous Contact Days_clean"]
for feature in continuous_features:
    plt.figure(figsize=(10,6))
    sns.boxplot(x='Subscription_Status_Code', y=feature, data=df_cleaned)
    plt.title(f'Boxplot of {feature} by Subscription Status')
    plt.xlabel('Subscription Status (0 = No, 1 = Yes)')
    plt.ylabel(feature)
    plt.show()
```











```
In [73]: cat_cols = ["Occupation", "Marital Status", "Education Level", "Credit Default",

for col in cat_cols:
    plt.figure(figsize=(12,6))
    sns.countplot(data=df_cleaned, x=col, hue="Subscription_Status_Code")
    plt.title(f"Distribution of {col} by Subscription Status")
    plt.xticks(rotation=45)
    plt.show()
```









(5) Feature Engineering

In order to make more meaningful insights, we will engineer some features such as:

- Job Type Category (Job_Type):
- Age Grouping (Age_Group):
- Education Type Category (Education_Rank):
- Loan Count (Loan_Count):

What is the Job_Type feature and how does it help the project?

Job_Type_Classification helps to group the original detailed occupations into a smaller number of meaningful categories.

They include:

White Collar: Contains people working in administration, management and entrepreneurs.

Blue Collar: Includes technicians, people who work in the services sector and blue collared workers

No Active Income: Consists of students, retirees and the unemployed, which are all groups that currently do not have an income, but may still have savings

Other: Contains unknown cases

White collar roles tend to usually earn more than blue collar roles.

Positions with a fixed salary are likely to be more financially stable than those who are unemployed or not working.

People with a higher income are likely to have a higher disposable income and may be more willing to commit their funds into a term deposit, which is a long term product.

This feature helps to decrease the wide spread of data across the many categories found in "Occupation", by combining them into broader but still meaningful groups. This helps

to increase the sample count while capturing the general importance of each category, and would lead to more reliable analysis and patterns.

```
In [74]: def Job_Type_Classification(job):
    if job in ["admin.", "management", "entrepreneur", "self-employed"]:
        return "White-Collar"
    elif job in ["blue-collar", "technician", "services", "housemaid"]:
        return "Blue-Collar"
    elif job in ["student", "unemployed", "retired"]:
        return "No Active Income"
    else:
        return "Other"

df_cleaned["Job_Type"] = df_cleaned["Occupation"].apply(Job_Type_Classification)
```

What is the Age_Group feature and how does it help the project?

Age_group helps to convert the numerical age into bins

Youth (<= 25 years old)

Young Adult (26-35 years old)

Adult (36-50 years old)

Mature (51-65 years old)

Senior (>65 years old)

Age is and has always been an extremely popular feature in the financial sector:

Younger people will generally tend to have lower savings, but they have This helps to improve the scope of the data and generalize the data into more meaningful variables. Also, based on research and other financial datasets I have looked at online, the relationship between age and the subscription is not strictly linear. Hence, binning helps to capture the relationship better by age groups than just singular numeric age values.

This feature helps to reduce noise and captures non-linear patterns better.

```
In [75]: bins=[0, 25, 35, 50, 65, 100]
labels=["Youth", "Young Adult", "Adult", "Mature", "Senior"]
df_cleaned["Age_Group"] = pd.cut(df_cleaned["Age_numerical"], bins=bins, labels=la
```

What is the Loan Count feature and how does it help the project?

Loan_Count aggregates the "Housing Loan" and "Personal Loan" into a singular integer, where 0 means no loans, 1 means one of the two loans, and 2 means both housing and personal loans.

This helps to improve the scope of the data and generalize the data into more meaningful variables.

Customers with 0 loans may have more disposable income and might be more willing to place their money in a term deposit.

Customers with 1 loan are likely to be in a moderate financial position.

Customers with 2 loans are likely to have less disposable income and might not be able

to afford to place their funds into a term deposit.

This feature helps to combine information from both "Housing Loan" and "Personal Loan", and could reveal patterns that could not be captured separately in the two columns.

```
In [76]: df_cleaned["Loan_Count"]=((df_cleaned["Housing Loan"]=="yes").astype(int) + (df_
```

(6) Label Encoding

For the final bivariate analysis, we will convert all the meaningful categorical columns into numerical, to improve the visualization.\

```
In [77]: for col in ["Job_Type", "Age_Group"]:
    if col not in cat_cols:
        cat_cols.append(col)
```

```
In [78]: df_cleaned[cat_cols].head()
```

Out[78]:

| | Occupation | Marital Status | Education Level | Credit Default | Housing Loan | Personal Loan | Contact Method | Job_ |
|---|-------------|----------------|---------------------|----------------|--------------|---------------|----------------|----------|
| 0 | technician | married | high.school | no | no | yes | cellular | C |
| 1 | blue-collar | married | basic.9y | no | no | no | cellular | C |
| 2 | admin. | married | high.school | no | no | no | telephone | W_C |
| 3 | housemaid | married | high.school | no | no_info | no | cellular | C |
| 4 | retired | married | professional.course | no | no_info | yes | cellular | No A_Inc |

```
In [79]: for col in cat_cols:
    df_cleaned[col + "_code"] = df_cleaned[col].astype('category').cat.codes
```

```
In [80]: for col in cat_cols:
    print(f"Mapping for {col}:")
    categories = df_cleaned[col].astype('category').cat.categories
    mapping_df = pd.DataFrame({col: categories,
                               col + 'Code': range(len(categories))})
    display(mapping_df)
```

Mapping for Occupation:

Occupation OccupationCode

| | | |
|-----------|---------------|----|
| 0 | admin. | 0 |
| 1 | blue-collar | 1 |
| 2 | entrepreneur | 2 |
| 3 | housemaid | 3 |
| 4 | management | 4 |
| 5 | retired | 5 |
| 6 | self-employed | 6 |
| 7 | services | 7 |
| 8 | student | 8 |
| 9 | technician | 9 |
| 10 | unemployed | 10 |

Mapping for Marital Status:

Marital Status Marital StatusCode

| | | |
|----------|----------|---|
| 0 | divorced | 0 |
| 1 | married | 1 |
| 2 | single | 2 |

Mapping for Education Level:

Education Level Education LevelCode

| | | |
|----------|---------------------|---|
| 0 | basic.4y | 0 |
| 1 | basic.6y | 1 |
| 2 | basic.9y | 2 |
| 3 | high.school | 3 |
| 4 | professional.course | 4 |
| 5 | university.degree | 5 |

Mapping for Credit Default:

Credit Default Credit DefaultCode

| | | |
|----------|---------|---|
| 0 | no | 0 |
| 1 | unknown | 1 |

Mapping for Housing Loan:

Housing Loan Housing LoanCode

| | | |
|----------|---------|---|
| 0 | no | 0 |
| 1 | no_info | 1 |
| 2 | yes | 2 |

Mapping for Personal Loan:

Personal Loan Personal LoanCode

| | | |
|----------|---------|---|
| 0 | no | 0 |
| 1 | no_info | 1 |
| 2 | yes | 2 |

Mapping for Contact Method:

Contact Method Contact MethodCode

| | | |
|----------|-----------|---|
| 0 | cellular | 0 |
| 1 | telephone | 1 |

Mapping for Job_Type:

Job_Type Job_TypeCode

| | | |
|----------|------------------|---|
| 0 | Blue-Collar | 0 |
| 1 | No Active Income | 1 |
| 2 | White-Collar | 2 |

Mapping for Age_Group:

Age_Group Age_GroupCode

| | | |
|----------|-------------|---|
| 0 | Youth | 0 |
| 1 | Young Adult | 1 |
| 2 | Adult | 2 |
| 3 | Mature | 3 |
| 4 | Senior | 4 |

In [81]: `df_cleaned.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38285 entries, 0 to 38284
Data columns (total 27 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Client ID        38285 non-null   int64  
 1   Occupation       38285 non-null   object  
 2   Marital Status   38285 non-null   object  
 3   Education Level  38285 non-null   object  
 4   Credit Default   38285 non-null   object  
 5   Housing Loan     38285 non-null   object  
 6   Personal Loan    38285 non-null   object  
 7   Contact Method   38285 non-null   object  
 8   Campaign Calls   38285 non-null   int64  
 9   Subscription Status 38285 non-null   object  
 10  Age_numerical   38285 non-null   int64  
 11  Age_Invalid     38285 non-null   int64  
 12  no_previous_contact 38285 non-null   int64  
 13  Previous Contact Days_clean 1366 non-null   float64 
 14  Subscription_Status_Code 38285 non-null   int64  
 15  Job_Type         38285 non-null   object  
 16  Age_Group        38285 non-null   category 
 17  Loan_Count       38285 non-null   int64  
 18  Occupation_code 38285 non-null   int8   
 19  Marital Status_code 38285 non-null   int8   
 20  Education Level_code 38285 non-null   int8   
 21  Credit Default_code 38285 non-null   int8   
 22  Housing Loan_code 38285 non-null   int8   
 23  Personal Loan_code 38285 non-null   int8   
 24  Contact Method_code 38285 non-null   int8   
 25  Job_Type_code   38285 non-null   int8   
 26  Age_Group_code   38285 non-null   int8  
dtypes: category(1), float64(1), int64(7), int8(9), object(9)
memory usage: 5.3+ MB

```

Final bivariate analysis takeaways


There are no features with high linearity, although it may appear that there are some. Those few that have above a 60% linearity is due to them being engineered features. We kept in the original features as well to see if there are any nuanced relationships

```
In [82]: numeric_cols= df_cleaned.select_dtypes(include=['number']).columns

numeric_cols= numeric_cols.drop(['Client ID'])

plt.figure(figsize=(16,12))
sns.heatmap(df_cleaned[numeric_cols].corr(numeric_only=True), annot=True)
```

Out[82]: <Axes: >



Important Insights from the charts

We can make a few important insights from the charts below, especially based on these few classes:

no_previous_contact ("0" category): For people who were previously contacted, there was a significantly higher chance of them subscribing to the term deposit, with a percentage of 63%. This could point to the fact that having contact could boost or already signify the fact that the client already has a strong interest and wants to gather more information on the terms and conditions of the term deposit.


Contact Method ("cellular" category): Clients who were last contacted with through cellular, had a slightly higher subscription rate of 14% compare to 4% by those contacted through telephone.

Age_Group ("Senior" category) & Age_numerical ("0" category): Looking at the Age_Group chart, we can see that there is a generally trend of Seniors (>65 age) having a higher subscription rate of over 40% and is double of the next highest category which is the Youth.


Job_Type ("No Active Income" category) & Occupation ("student", "retired", "unemployed" category): Generally, the graphs show how people with no active income

are having a higher subscription rate that is about twice the amount of other job types. Looking deeper into the specific occupations, we can see how students have the highest rate followed by the retired and the unemployed.\

```
In [83]: age_group_sub = df_cleaned.groupby("Age_Group", observed=True)[ "Subscription_Status"]  
  
plt.figure(figsize=(10,6))  
sns.barplot(x=age_group_sub.index, y=age_group_sub.values)  
plt.title("Average Subscription Rate by Age Group")  
plt.show()
```




```
In [84]: plt.figure(figsize=(12,6))  
sns.histplot(data=df_cleaned, x="Age_numerical", hue="Subscription_Status_Code",  
plt.title("Age Distribution by Subscription Status")  
plt.xlabel("Age")  
plt.ylabel("Density")  
plt.legend(title="Subscription Status", labels=["No", "Yes"])  
plt.show()
```




```
In [85]: job_type_sub= df_cleaned.groupby("Job_Type")["Subscription_Status_Code"].mean()

plt.figure(figsize=(10,6))
sns.barplot(x=job_type_sub.index, y=job_type_sub.values)
plt.title("Average Subscription Rate by Job Type")
plt.show()
```




```
In [86]: occupation_sub= df_cleaned.groupby("Occupation")["Subscription_Status_Code"].mean()
plt.figure(figsize=(14,8))
sns.barplot(x=occupation_sub.index, y=occupation_sub.values)
plt.title("Average Subscription Rate by Occupation")
plt.show()
```




```
In [87]: loan_count_sub= df_cleaned.groupby("Loan_Count")["Subscription_Status_Code"].mean()

plt.figure(figsize=(10,6))
sns.barplot(x=loan_count_sub.index, y=loan_count_sub.values)
plt.title("Average Subscription Rate by Loan Count")
plt.show()
```




```
In [88]: contact_method_sub= df_cleaned.groupby("Contact Method")["Subscription_Status_Code"].mean()


plt.figure(figsize=(10,6))
sns.barplot(x=contact_method_sub.index, y=contact_method_sub.values)
plt.title("Average Subscription Rate by Contact Method")
plt.show()
```




```
In [89]: education_level_sub= df_cleaned.groupby("Education Level")["Subscription_Status_Co  
plt.figure(figsize=(10,6))  
sns.barplot(x=education_level_sub.index, y=education_level_sub.values)  
plt.title("Average Subscription Rate by Education Level")  
plt.show()
```




```
In [90]: marital_status_sub= df_cleaned.groupby("Marital Status")["Subscription_Status_Co  
plt.figure(figsize=(10,6))  
sns.barplot(x=marital_status_sub.index, y=marital_status_sub.values)  
plt.title("Average Subscription Rate by Marital Status")  
plt.show()
```



```
In [91]: prev_contact_sub= df_cleaned.groupby("no_previous_contact")["Subscription_Status"].mean()
plt.figure(figsize=(10,6))
sns.barplot(x=prev_contact_sub.index, y=prev_contact_sub.values)
plt.title("Average Subscription Rate by Previous Contact Status")
plt.show()
```



```
In [92]: calls_sub= df_cleaned.groupby("Campaign Calls")["Subscription_Status_Code"].mean()
plt.figure(figsize=(12,6))
sns.barplot(x=calls_sub.index, y=calls_sub.values)
plt.title("Number of Campaign Calls to Subscription ratio")
plt.show()
```



(7) Key Insights Summary

Based on the exploratory data analysis conducted on this bank marketing dataset, we have observed many important patterns based on the clients' behaviours and different demographics. These insights we have concluded, can definitely help the bank refine their future marketing strategies, as well as which customer segments have high potential to be subscribers to the term deposits.

1. The overall subscription rate is very low
 - Only ~11% of the clients subscribed to the term deposit.
 - The dataset is extremely imbalanced, with almost 9/10 clients not going through with the subscription.
 - This could suggest how the mass-marketing approaches done by the company are not efficient and are expensive .
2. Having a previous contact history significantly boosts subscription likelihood
 - Clients who had been previously contacted showed a ~63% subscription rate, which is very good.
 - Extremely different from first-time contacts which constitutes to ~9 of the subscription rate.
 - This has been shown to be the predictor with the highest impact out of all the features we have analysed.
3. Age-Subscription Trend
 - The subscription rate varies and it resembles a "U" shape, with Seniors having 45% of the distribution and the Youth having 19% out of all the people who subscribed.
 - This shows how seniors show the highest interest. Followed by the Youth and then the Young Adults and Mature clients.
 - The Adult age group has the lowest subscription rate, and this could indicate possible financial commitments such as saving for a house, or paying off their car, or even raising a family.

-Hence, this U-shaped pattern helps to highlight that the youngest and oldest age groups are the most receptive and responsive to the term deposit offer

4. Job type influences the subscription rate

-No Active Income clients have a higher subscription rate than the rest, at ~23%.

-This also aligns with the Age Group results as seniors and students are more likely to subscribe.

-White-collar clients have a higher subscription rate than blue-collar clients.

-This could suggest how financial stability and the current available savings can contribute to the client's decision to subscribe.

5. Contact method's effect on subscription rates

-Although not extreme, cellular methods of contacting the clients have shown to perform better than the traditional telephone method.

-Moving forward, a digital-first method will likely prove to be a more viable strategy.

6. Education Level scales with subscription rates

-There is a general trend that the higher the education level of a client, the more likely it is for the client to subscribe to the term deposit. This shows that there is a linear relationship between the two features.

-This in fact proves how the higher the education level, the more knowledge the client has about finances and they tend to be more financially literate.

7. Loan Count and Marital Status are not major factors in influencing the client's decision

-We have observed how the number of loans and the marital status of the client do not have a noticeable effect on whether the client decides to subscribe or not.